

Securing Shared-Kernel Runtimes Against Kernel-Level Container Escapes

Hexa Force Core Engineering Team
DevSecOps and Cloud Security Research Group
Indian Institute of Technology Jodhpur (IITJ)
Jodhpur, India
contact@hexaforce.org

Abstract—Standard OS-level containerization technologies achieve lightweight performance and scalability by sharing the underlying host operating system kernel directly. However, this shared-kernel architectural model introduces a critical single point of failure: any local privilege escalation (LPE) vulnerability in the host kernel is immediately exposed to all container namespaces running on that host. This paper presents and analyzes *Hexa Force*, a multi-layered, zero-trust container security framework designed to detect, contain, and remediate container escapes originating from kernel-level vulnerabilities. We model a high-fidelity 4-stage exploit chain—starting from an unprivileged container footprint, leveraging the Dirty COW (CVE-2016-5195) virtual memory race condition to gain container root, abusing Linux administrative capabilities (CAP_SYS_ADMIN) and exposed Unix domain API sockets (/var/run/docker.sock) to escape, and executing a persistent root shell takeover of the host operating system. To counter this threat, we implement a comprehensive hardening strategy combining system-level system call sandboxing via custom Secure Computing Mode (Seccomp) profiles to block madvise system calls, process-level capability dropping, and state-of-the-art software supply chain integrity using Sigstore/Cosign and OpenID Connect (OIDC) keyless container signing inside automated GitHub Actions pipelines. Our experimental evaluation demonstrates a 100% exploit containment rate under active Seccomp profiles, zero false positives, and minimal deployment overhead, delivering an enterprise-ready blueprint for securing modern shared-kernel cloud-native workloads.

Index Terms—Container Security, DevSecOps, Seccomp System Call Filtering, Sigstore Cosign, OpenID Connect (OIDC), Kernel Privilege Escalation, CVE-2016-5195, Host Escape.

I. INTRODUCTION

Modern cloud-native software development relies heavily on containerization to achieve rapid continuous integration and continuous deployment (CI/CD). Unlike heavy hardware-level virtual machines (VMs) governed by a hypervisor, standard container runtimes (such as Docker, containerd, and CRI-O) share the physical host system's Linux kernel directly. This shared-kernel architectural model provides lightweight virtualization, high performance, and rapid start-up times.

However, sharing a kernel introduces severe security risks. Because container boundaries are logical constructs enforced by kernel namespaces, cgroups, and capabilities, any flaw inside the host's virtual memory subsystem or system call interface can collapse the entire isolation sandbox. If an attacker gains an unprivileged shell inside a container, they

can compile and execute a local privilege escalation (LPE) exploit that directly targets the unpatched host kernel.

This paper introduces the *Hexa Force* security architecture, demonstrating a rigorous attack-defense lifecycle. We model a comprehensive 4-stage exploit chain that demonstrates how a minor kernel race condition, paired with common runtime misconfigurations, can allow a low-privileged containerized application user to completely break isolation boundaries and take persistent root control of the physical host machine.

To counter this critical threat, we design a multi-layered, zero-trust DevSecOps defense blueprint. At the host-kernel boundary, we implement system-level system call filtering using a custom **Seccomp** profile that blocks the specific system calls used in memory race exploits, neutralizing the vulnerability at the kernel interface. At the container runtime level, we strip process privileges and capabilities. At the software supply-chain level, we integrate automated vulnerability scanning (using AquaSec Trivy) and cryptographic image signing (using keyless Sigstore/Cosign federated via GitHub Actions OIDC tokens) to ensure that only verified, signed, and scanned container images can ever enter the execution environment.

Our primary contributions are:

- A high-fidelity 4-stage threat emulation mapping container escapes from initial footprint to total host takeover.
- A mathematical and state-machine analysis of the CVE-2016-5195 virtual memory Copy-on-Write race condition.
- A concrete system call sandboxing implementation using a custom Seccomp profile to block memory-manipulation syscalls.
- An automated DevSecOps CI/CD pipeline using Trivy vulnerability layer checks and Cosign OIDC keyless cryptographic signatures.

—

II. THE 4-STAGE CONTAINER THREAT EMULATION MODEL

To analyze how logical container boundaries collapse under unpatched kernels, we designed a high-fidelity 4-stage exploit chain that models a realistic, cascading lateral-movement compromise.

A. Stage 1: Host Kernel Isolation (Dirty COW Exploit)

The threat begins inside an Ubuntu-based container running as a low-privileged user `victim`. This simulates a web application compromise (e.g., Remote Code Execution). The attacker compiles and runs `dirtycow.c` inside the container. This binary targets the Dirty COW kernel vulnerability (CVE-2016-5195), a race condition in the kernel’s virtual memory subsystem. By writing to a read-only target file (`/var/secret/target.txt`) using the `madvise` system call and the `/proc/self/mem` interface, the exploit bypasses standard page cache permissions. The attacker successfully overwrites the root-owned target, escalating their privilege from the unprivileged `victim` shell to `root` inside the container namespace.

B. Stage 2: Process Boundaries (Capability Abuse)

Obtaining container root is necessary, but namespaces still limit the attacker’s visibility. The attacker now queries the process’s Linux capabilities. In standard Docker configurations, containers are sometimes granted excessive administrative privileges. If the container is misconfigured with the `CAP_SYS_ADMIN` capability (often done to support nested containers or legacy monitoring agents), the logical boundaries of namespace isolation crumble. The attacker leverages `CAP_SYS_ADMIN` to configure filesystem mount layers, preparing the physical groundwork for escaping the container namespace.

C. Stage 3: Engine API Security (docker.sock Hijack)

A common development practice is mounting the host’s physical Docker socket `/var/run/docker.sock` directly inside the container namespace to enable continuous integration runners or administrative agents. Because the attacker obtained container root in Stage 1, they have full read and write permissions to this mounted socket. The Docker socket is a Unix domain socket that connects directly to the host’s root-privileged Docker daemon. The attacker hijacks this socket, issuing direct HTTP REST API calls to command the host Docker engine to pull external images and instantiate new privileged container configurations.

D. Stage 4: Volume & Filesystem Security (Host Takeover)

Using their control over the Docker socket, the attacker commands the host engine to launch a new, highly privileged container. In this container’s deployment configuration, the attacker mounts the physical host’s entire root filesystem `/` to `/mnt/host` inside the container. Since the container is privileged, the attacker can traverse directly into the host’s directory structure. They navigate into `/mnt/host/etc/cron.d` and write a malicious cron job. When the host operating system’s cron daemon periodically executes this script, it triggers a root-privileged reverse shell back to the attacker’s system. This establishes a persistent root backdoor on the physical host machine, completing a full container breakout.

III. MATHEMATICAL & OS-LEVEL VULNERABILITY ANALYSIS

The core of our threat model relies on the race condition in the virtual memory subsystem’s handling of **Copy-on-Write (COW)** during private read-only file mappings (`MAP_PRIVATE`).

Let P_{shared} represent a physical memory page cache frame corresponding to a read-only backing file mapped privately into process virtual memory space V . Under normal kernel operations, when a process attempts to execute a write operation against V , the CPU’s Memory Management Unit (MMU) throws a write page fault because the page table entry is marked read-only.

The kernel page fault handler intercepts this, allocating a new, private physical memory page frame P_{private} , copying the contents of P_{shared} to P_{private} , updating the process’s page table entry to link $V \rightarrow P_{\text{private}}$, and completing the write. This is the Copy-on-Write process.

The Dirty COW exploit (CVE-2016-5195) targets a **non-atomic race condition** inside this page fault handler. The exploit spawns two concurrent, high-speed threads:

A. The Discard Thread (Thread A)

This thread runs in a tight loop calling `madvise` with the `MADV_DONTNEED` flag:

$$\text{madvise}(V, \text{size}, \text{MADV_DONTNEED}) \quad (1)$$

This system call advises the kernel that the process no longer needs the specified virtual address range. The kernel responds by clearing the physical page table pointers mapping $V \rightarrow P_{\text{private}}$, discarding the private copy.

B. The Write Thread (Thread B)

Concurrently, this thread writes the payload to the virtual process memory file `/proc/self/mem` targeting virtual address V :

$$\text{lseek}(fd_mem, V, \text{SEEK_SET}); \quad \text{write}(fd_mem, \text{payload}, \text{size}) \quad (2)$$

Writing to `/proc/self/mem` routes the write operation directly through the kernel’s virtual memory handler, bypassing standard MMU write permission checks and triggering the vulnerable page fault handler.

C. The Race Condition Collision

Because the page fault routine is non-atomic, the following execution sequence occurs across separate CPU cores:

- 1) **Thread B** triggers a write, causing the kernel to check write permissions. The check succeeds (COW is required).
- 2) The kernel initiates the COW allocation, allocating and copying the contents to a private page.
- 3) **Thread A** executes `madvise(MADV_DONTNEED)`, discarding the virtual page table entry.
- 4) The kernel, resuming the interrupted write operation of **Thread B**, attempts to update the page tables. However,

because the page was discarded, the kernel gets confused and falls back to writing directly to the original, read-only backing physical page frame P_{shared} in the page cache.

- 5) The kernel executes the physical write, modifying the underlying file on disk.

D. The Kernel Patch Analysis

The Linux kernel community resolved this vulnerability by introducing an internal state flag called `FOLL_COW` within the virtual memory follow-page handlers. The patched code guarantees that the kernel tracks whether a Copy-on-Write operation has already been initiated for the page. When `madvise` concurrently discards the page table mapping, the `FOLL_COW` flag serves as an internal “sticky bit,” ensuring the kernel remembers that it is in the middle of a private COW allocation, preventing it from ever falling back to writing directly to the shared read-only page cache.

IV. PROACTIVE SYSTEM-LEVEL HARDENING

To protect against kernel-level virtual memory exploits on shared-kernel runtimes, security administrators must implement proactive system call filtering at the container-to-host boundary.

We designed a custom, production-grade **Seccomp (Secure Computing Mode)** security profile named `hexaforce-seccomp.json`. Seccomp is a Linux kernel feature that allows system administrators to restrict the system calls a containerized process can execute, drastically reducing the kernel attack surface.

```
{
  "defaultAction": "SCMP_ACT_ALLOW",
  "architectures": [
    "SCMP_ARCH_X86_64",
    "SCMP_ARCH_X86"
  ],
  "syscalls": [
    {
      "names": [ "madvise" ],
      "action": "SCMP_ACT_ERRNO",
      "comment": "Block madvise system calls."
    }
  ]
}
```

By default, Docker’s default Seccomp profile allows the `madvise` system call, as it is a common memory-discarding optimization call used by memory managers. However, standard microservices (such as web APIs or databases) have no legitimate operational need to call `madvise` at runtime.

By executing the container using the custom Hexa Force Seccomp profile:

```
docker run --rm -it --security-opt
  seccomp=hexaforce-seccomp.json
  dirtycow-lab
```

the kernel instantly intercepts any call to `madvise`, immediately returning an `EPERM` (Operation not permitted) error code to the process. When the `dirtycow` binary launches inside the container, Thread A’s loop fails on the very first iteration, completely neutralizing the memory page discard mechanism and preventing the race condition from ever occurring.

V. CRYPTOGRAPHIC SUPPLY-CHAIN TRUST

Securing the runtime environment is only half the battle; we must also guarantee **code provenance** to ensure that unauthorized, modified, or vulnerable container images can never enter our production execution clusters. Hexa Force implements state-of-the-art keyless container signing using the **Sigstore/Cosign** ecosystem federated via **OpenID Connect (OIDC)**.

A. The Keyless Signing Architecture

Traditional container signing depends on managing cryptographic private keys stored inside CI/CD environment variables. If these private keys are leaked or stolen, the integrity of the entire software supply chain collapses.

Hexa Force eliminates this risk entirely by using **keyless signing**:

- 1) Upon a code push, our **GitHub Actions** runner requests an ephemeral identity token from the trusted GitHub OIDC provider.
- 2) The runner sends this OIDC token to the **Fulcio Certificate Authority** (a Sigstore service).
- 3) Fulcio verifies the token’s cryptographic signatures, confirms the build pipeline’s identity (e.g., repository URL, commit hash, workflow path), and issues a short-lived (10-minute) X.509 signing certificate.
- 4) **Cosign** uses this ephemeral certificate to sign the container image and pushes the signature metadata to the public **Rekor immutable transparency ledger**.
- 5) The private key is instantly discarded by Fulcio, completely removing the threat of private key leaks.

B. Automated Pipeline Gates

We fully automated this supply chain trust model inside a robust GitHub Actions workflow (`dirtycow-lab.yml`). The workflow enforces the following gates:

- 1) **Source Linting & Build**: Compiles the source files and builds the container image locally.
- 2) **Static Vulnerability Scanning**: Executes **AquaSec Trivy** to scan all image layers for `CRITICAL` or `HIGH` vulnerabilities, reporting any security regressions in the GHA logs.
- 3) **Registry Push**: Dynamically converts the repository tag to lowercase (complying with OCI standards) and pushes the verified image to the **GitHub Container Registry (GHCR)**.

- 4) **OIDC Cryptographic Signature:** Automatically executes the Cosign keyless signing step using the runner’s ephemeral Fulcio identity.
-

VI. EXPERIMENTAL RESULTS & PERFORMANCE EVALUATION

To validate the effectiveness and performance characteristics of our multi-layered defense architecture, we conducted rigorous comparative experiments.

A. Exploit Containment Rates

We compared the success rate of the Dirty COW exploit across three distinct environment configurations. The experiment executed 100 iterations of the exploit binary per configuration:

- **Configuration 1 (Default / Vulnerable):** The container was run with default capabilities. The Dirty COW exploit succeeded in **100% of runs**, overwriting `/var/secret/target.txt` within an average of 3.2 seconds.
- **Configuration 2 (Capability Dropping Only):** Dropping `CAP_SYS_ADMIN` blocked the Stage 2 namespace escape. However, because `madvise` was still allowed, the exploit successfully escalated the process privilege to container-root in **100% of runs**.
- **Configuration 3 (Custom Seccomp Profile):** With the custom Seccomp profile active, the exploit failed in **100% of runs**, resulting in an immediate exit code and zero page cache modifications, verifying absolute containment.

B. Deployment & Execution Overhead

We evaluated the performance overhead introduced by our security controls. Restricting Linux capabilities and applying custom Seccomp filters are logical security policies enforced natively at the kernel syscall boundary.

Our benchmark results verified that Seccomp filtering introduces **0.00% physical CPU and memory overhead** during runtime application execution, making it highly suitable for high-performance enterprise workloads.

Furthermore, our GHA pipeline’s automated Trivy scanning layer required an average of only **14 seconds** to execute, providing robust DevSecOps protection with negligible impact on pipeline delivery speed.

VII. RELATED WORK & FUTURE EVOLUTIONARY PATHS

While our multi-layered hardening strategy provides exceptionally strong, production-grade security, the scientific community recognizes that namespace-based isolation still shares a core vulnerability: the host kernel.

A. Micro-VM Runtimes

To isolate the kernel boundary completely, advanced enterprises can transition from standard shared-kernel runtimes to **user-space sandboxing or micro-virtualization**:

- 1) **gVisor (Google):** Introduces a user-space guest kernel (written in Go) that intercepts all system calls executed by the container, handling them inside the sandbox and completely blocking direct system calls to the host kernel.
 - 2) **Kata Containers:** Executes each container inside a dedicated, ultra-lightweight micro-virtual machine (using QEMU or Firecracker) with its own guest kernel, achieving absolute, physical hardware-level isolation.
-

VIII. CONCLUSION & STRATEGIC OUTLOOK

This paper successfully demonstrates the fragility of default shared-kernel container environments and delivers a robust, multi-layered DevSecOps architecture designed to contain and neutralize kernel-level exploits. By combining proactive system call filtering (via custom Seccomp profiles) with automated vulnerability layer scanning (Trivy) and cryptographic software supply-chain integrity validation (Sigstore Cosign and OpenID Connect), the **Hexa Force** blueprint achieves complete vulnerability containment with zero operational overhead.

Our research proves that container security must be treated as a continuous lifecycle rather than a static logical sandbox, and our DevSecOps integration model delivers a practical, enterprise-ready standard for modern zero-trust cloud infrastructure.

REFERENCES

- [1] P. Oester, “CVE-2016-5195: Linux Kernel Copy-on-Write Privilege Escalation Vulnerability,” MITRE Common Vulnerabilities and Exposures, Oct. 2016.
- [2] Sigstore Community, “Sigstore: Cryptographic Software Supply Chain Integrity,” Linux Foundation Research, 2021. [Online]. Available: <https://sigstore.dev>
- [3] Docker Inc., “Docker Security and Seccomp System Call Filtering Profiles,” Docker Documentation, 2023. [Online]. Available: <https://docs.docker.com/engine/security/seccomp/>
- [4] Aquasecurity, “Trivy: High-performance Container Vulnerability Scanner,” GitHub Repository, 2023. [Online]. Available: <https://github.com/aquasecurity/trivy>
- [5] Google Cloud, “gVisor: Sandbox Container Runtime Interface,” GitHub Repository, 2020. [Online]. Available: <https://github.com/google/gvisor>
- [6] Open Container Initiative (OCI), “OCI Runtime Specification and Image Standards,” Linux Foundation, 2022.